

Titanic_EDA_Project

Titanic_EDA_Project.ipynb

```
import pandas as pd
df = pd.read_csv("/content/train (1).csv")
display(df.head())
```


| | PassengerId | Survived | Pclass | Name | Sex | Age | SibSp | Parch | Ticket | Fare | Cabin | Embarked |
|---|-------------|----------|--------|--|--------|------|-------|-------|---------------------|---------|-------|----------|
| 0 | 1 | 0 | 3 | Braund, Mr. Owen Harris | male | 22.0 | 1 | 0 | A/5 21171 | 7.2500 | NaN | S |
| 1 | 2 | 1 | 1 | Cumings, Mrs. John Bradley (Florence Briggs Th... | female | 38.0 | 1 | 0 | PC 17599 | 71.2833 | C85 | C |
| 2 | 3 | 1 | 3 | Heikinen, Miss. Laina | female | 26.0 | 0 | 0 | STON/O2. 3101282 | 7.9250 | NaN | S |
| 3 | 4 | 1 | 1 | Futrelle, Mrs. Jacques Heath (Lily | female | 35.0 | 1 | 0 | 113803 | 53.1000 | C123 | S |

```
# Import Visualization Libraries
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style('whitegrid')
```

Survival Count

Explanation: This plot shows the number of passengers who survived vs. those who did not. It helps us see the overall survival rate on the Titanic.

```
# Plot Survival Count
sns.countplot(x='Survived', data=df)
plt.title("Survival Count (0 = Died, 1 = Survived)")
plt.show()
```




Observation: A significantly higher number of passengers did not survive compared to those who did, highlighting the tragedy of the Titanic disaster.

Gender Distribution

Explanation: This plot shows the number of male and female passengers on the Titanic. It helps us understand the gender balance in the dataset.

```
# Plot Gender Distribution
sns.countplot(x='Sex', data=df)
```

```
plt.title("Gender Distribution")  
plt.show()
```




Observation: There were significantly more male passengers than female passengers on board the Titanic.

✎ Passenger Class Distribution

Explanation: This plot shows how many passengers were in each class (1st, 2nd, 3rd). It gives an idea of the Titanic's class structure and how passengers were distributed.

```
# Plot Passenger Class Distribution  
sns.countplot(x='Pclass', data=df)  
plt.title("Passenger Class Distribution")  
plt.show()
```




Observation: Most passengers were in 3rd class, followed by 1st and 2nd. This shows that the Titanic carried a large number of lower-class passengers.

✎ Survival by Gender

Explanation: This plot shows survival rates grouped by gender. It helps us understand whether men or women had higher chances of survival on the Titanic.

```
# Survival by Gender  
sns.countplot(x='Sex', hue='Survived', data=df)
```

```
plt.title("Survival by Gender")  
plt.show()
```




Observation: A much higher proportion of women survived compared to men. This suggests that gender played a major role in survival – possibly due to the "women and children first" evacuation policy.

✓ Survival by Passenger Class

Explanation: This plot shows survival counts across different passenger classes (1st, 2nd, 3rd). It helps us understand whether class had an impact on survival rates.

```
# Survival by Passenger Class  
sns.countplot(x='Pclass', hue='Survived', data=df)  
plt.title("Survival by Passenger Class")  
plt.show()
```




Observation: Passengers in 1st class had a much higher survival rate than those in 2nd or 3rd class. Survival chances dropped significantly for 3rd class passengers.

✓ Boxplot: Fare vs Survival

Explanation: This boxplot compares the distribution of fares for passengers who survived and those who didn't. It helps reveal if paying higher fares increased the chance of survival.

```
# Boxplot: Fare vs. Survival
sns.boxplot(x='Survived', y='Fare', data=df)
plt.title("Fare Amount vs. Survival")
plt.show()
```




Observation: Passengers who paid higher fares had a higher chance of survival. This indicates that people in higher classes (who paid more) were more likely to survive.

Boxplot: Age vs Survival

Explanation: This boxplot displays the age distribution for passengers who survived vs those who didn't. It can reveal if younger passengers had better survival chances.

```
sns.boxplot(x='Survived', y='Age', data=df)
plt.title("Boxplot: Age vs Survival")
plt.xlabel("Survived (0 = No, 1 = Yes)")
plt.ylabel("Age")
plt.show()
```




Observation: The age distribution of survivors shows a mix, but there may be a slightly higher concentration of younger survivors. Some children and young adults had better chances.

✓ Checking for Missing Values

We identify columns with missing data before handling them.

```
df.isnull().sum()
```



| | 0 |
|--------------------|-----|
| PassengerId | 0 |
| Survived | 0 |
| Pclass | 0 |
| Name | 0 |
| Sex | 0 |
| Age | 177 |
| SibSp | 0 |
| Parch | 0 |
| Ticket | 0 |
| Fare | 0 |
| Cabin | 687 |
| Embarked | 2 |


dtype: int64

✓ Data Preprocessing (Clean Start)

Dropping Irrelevant Columns

Removed columns that do not contribute meaningfully to survival prediction.

```
df = df.drop(columns=[col for col in ['Name', 'Ticket', 'Cabin'] if col in df.columns])
print(df.columns)
```



```
Index(['PassengerId', 'Survived', 'Pclass', 'Sex', 'Age', 'SibSp', 'Parch',
      'Fare', 'Embarked'],
      dtype='object')
```

We identify columns with missing data before handling them.

✓ Filling Missing Values

- Age filled with median
- Embarked filled with most frequent value

```
df['Age'] = df['Age'].fillna(df['Age'].median())
```

✓ Encoding Categorical Columns

- 'Sex' and 'Embarked' converted into numerical labels for modeling.

```
from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()
df['Sex'] = le.fit_transform(df['Sex'])
```

```
df.head()
```



| | PassengerId | Survived | Pclass | Sex | Age | SibSp | Parch | Fare | Embarked |
|---|-------------|----------|--------|-----|------|-------|-------|---------|----------|
| 0 | 1 | 0 | 3 | 1 | 22.0 | 1 | 0 | 7.2500 | S |
| 1 | 2 | 1 | 1 | 0 | 38.0 | 1 | 0 | 71.2833 | C |
| 2 | 3 | 1 | 3 | 0 | 26.0 | 0 | 0 | 7.9250 | S |
| 3 | 4 | 1 | 1 | 0 | 35.0 | 1 | 0 | 53.1000 | S |
| 4 | 5 | 0 | 3 | 1 | 35.0 | 0 | 0 | 8.0500 | S |

✓ Splitting Dataset

- 80% for training, 20% for testing
- random_state=42 ensures reproducibility

```
df.select_dtypes(include='object').columns
```



```
Index(['Embarked'], dtype='object')
```

```
# Drop text columns that are not useful for model
df = df.drop(columns=[col for col in ['Name', 'Ticket', 'Cabin'] if col in df.columns])
```

```
# Encode 'Sex' and 'Embarked' if they're still strings
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
```

```
if df['Sex'].dtype == 'object':
    df['Sex'] = le.fit_transform(df['Sex'])
```

```
if 'Embarked' in df.columns and df['Embarked'].dtype == 'object':
    df['Embarked'] = le.fit_transform(df['Embarked'])
```

```
from sklearn.model_selection import train_test_split
```

```
X = df.drop('Survived', axis=1)
y = df['Survived']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Train models
# a.Logistic Regression
from sklearn.linear_model import LogisticRegression
lr = LogisticRegression(max_iter=200)
lr.fit(X_train, y_train)
```

```
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
# Prediction
y_pred_lr = lr.predict(X_test)
```

```
# Evaluation performance
print("Logistic Regression Accuracy:", accuracy_score(y_test, y_pred_lr))
print("\nClassification Report:\n", classification_report(y_test, y_pred_lr))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred_lr))
```

```
# b.Decision Tree
from sklearn.tree import DecisionTreeClassifier
dt = DecisionTreeClassifier()
dt.fit(X_train, y_train)
```

```
# Predictions
y_pred_dt = dt.predict(X_test)
```

```
# Evaluation performance
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
print("Decision Tree Accuracy:", accuracy_score(y_test, y_pred_dt))
print("\nClassification Report:\n", classification_report(y_test, y_pred_dt))
```

```
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred_dt))
```

```
# c.Random Forest
from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier()
rf.fit(X_train, y_train)

# Predictions
y_pred_rf = rf.predict(X_test)

# Evaluation
print("Random Forest Accuracy:", accuracy_score(y_test, y_pred_rf))
print("\nClassification Report:\n", classification_report(y_test, y_pred_rf))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred_rf))
```

 Logistic Regression Accuracy: 0.8100558659217877

```
Classification Report:
              precision    recall  f1-score   support

     0       0.83         0.86         0.84         105
     1       0.79         0.74         0.76          74

   accuracy          0.81
  macro avg       0.81         0.80         0.80         179
 weighted avg       0.81         0.81         0.81         179
```

```
Confusion Matrix:
[[90 15]
 [19 55]]
Decision Tree Accuracy: 0.7597765363128491
```

```
Classification Report:
              precision    recall  f1-score   support

     0       0.79         0.80         0.80         105
     1       0.71         0.70         0.71          74

   accuracy          0.76
  macro avg       0.75         0.75         0.75         179
 weighted avg       0.76         0.76         0.76         179
```

```
Confusion Matrix:
[[84 21]
 [22 52]]
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
 Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
 n_iter_i = _check_optimize_result(
 Random Forest Accuracy: 0.8212290502793296

```
Classification Report:
              precision    recall  f1-score   support

     0       0.82         0.89         0.85         105
     1       0.82         0.73         0.77          74

   accuracy          0.82
  macro avg       0.82         0.81         0.81         179
 weighted avg       0.82         0.82         0.82         179
```

```
Confusion Matrix:
[[93 12]
 [20 54]]
```

Support Vector Machine (SVM)

```
# Support Vector Machine (SVM)

from sklearn.svm import SVC
```

```

svm = SVC()
svm.fit(X_train, y_train)

# Predict and evaluate
y_pred_svm = svm.predict(X_test)

print("SVM Accuracy:", accuracy_score(y_test, y_pred_svm))
print("\nClassification Report:\n", classification_report(y_test, y_pred_svm))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred_svm))

```

↔ SVM Accuracy: 0.5977653631284916

| Classification Report: | | precision | recall | f1-score | support |
|------------------------|------|-----------|--------|----------|---------|
| 0 | 0.60 | 0.98 | 0.74 | 105 | |
| 1 | 0.67 | 0.05 | 0.10 | 74 | |
| accuracy | | | 0.60 | 179 | |
| macro avg | 0.63 | 0.52 | 0.42 | 179 | |
| weighted avg | 0.62 | 0.60 | 0.48 | 179 | |

Confusion Matrix:

```
[[103  2]
 [ 70  4]]
```

K-Nearest Neighbors (KNN)

```

# K-Nearest Neighbours(KNN)
from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier()
knn.fit(X_train, y_train)

# Predict and evaluate
y_pred_knn = knn.predict(X_test)

print("KNN Accuracy:", accuracy_score(y_test, y_pred_knn))
print("\nClassification Report:\n", classification_report(y_test, y_pred_knn))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred_knn))

```

↔ KNN Accuracy: 0.659217877094972

| Classification Report: | | precision | recall | f1-score | support |
|------------------------|------|-----------|--------|----------|---------|
| 0 | 0.67 | 0.84 | 0.74 | 105 | |
| 1 | 0.64 | 0.41 | 0.50 | 74 | |
| accuracy | | | 0.66 | 179 | |
| macro avg | 0.65 | 0.62 | 0.62 | 179 | |
| weighted avg | 0.65 | 0.66 | 0.64 | 179 | |

Confusion Matrix:

```
[[88 17]
 [44 30]]
```

Naive Bayes

```

# Naive Bayes
from sklearn.naive_bayes import GaussianNB

nb = GaussianNB()
nb.fit(X_train, y_train)

# Predict and evaluate
y_pred_nb = nb.predict(X_test)

print("Naive Bayes Accuracy:", accuracy_score(y_test, y_pred_nb))
print("\nClassification Report:\n", classification_report(y_test, y_pred_nb))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred_nb))

```


Naive Bayes Accuracy: 0.776536312849162

```

Classification Report:
              precision    recall  f1-score   support

     0       0.83         0.78         0.80         105
     1       0.71         0.77         0.74          74

   accuracy          0.78         0.78         0.78         179
  macro avg          0.77         0.78         0.77         179
 weighted avg          0.78         0.78         0.78         179

```

```

Confusion Matrix:
[[82 23]
 [17 57]]

```

Model Comparison

```
# Model Comparison and Selection
```

```

models = {
    "Logistic Regression": accuracy_score(y_test, y_pred_lr),
    "Decision Tree": accuracy_score(y_test, y_pred_dt),
    "Random Forest": accuracy_score(y_test, y_pred_rf),
    "SVM": accuracy_score(y_test, y_pred_svm),
    "KNN": accuracy_score(y_test, y_pred_knn),
    "Naive Bayes": accuracy_score(y_test, y_pred_nb)
}

```

```

# Display model accuracies
for model_name, acc in models.items():
    print(f"{model_name}: {acc:.4f}")

```

Logistic Regression: 0.8101
 Decision Tree: 0.7598
 Random Forest: 0.8212
 SVM: 0.5978
 KNN: 0.6592
 Naive Bayes: 0.7765

Visualize Comparison

```
# Visualize Comparison
```


```

import matplotlib.pyplot as plt

model_names = list(models.keys())
accuracy_scores = list(models.values())

plt.figure(figsize=(10, 6))
plt.bar(model_names, accuracy_scores, color='skyblue')
plt.xlabel('Model')
plt.ylabel('Accuracy')
plt.title('Model Accuracy Comparison')
plt.xlim(0.5, 1)
plt.grid(axis='x')
plt.show()

```



Best Model

```
best_model = rf # assuming Random Forest performed best
```

```
import joblib
```

```
joblib.dump(best_model, 'titanic_model.pkl') # saves your model to file
```

```
model = joblib.load("titanic_model.pkl")  
prediction = model.predict(X_test)
```